

Course: DVA262 — Machine Learning Concepts

Objective: The purpose of the Lab is to apply your understanding of the k-NN algorithm by implementing the k-NN algorithm from scratch in the C++ programming using the Lab skeleton.

Lab 1: The k-NN Algorithm Implementation

Execution and demonstration

This lab is executed by a group of two students or individually. The completed lab should be demonstrated to a lab assistant.

Useful aids

- The course literature
- Lecture slides and notes
- Eigen <https://eigen.tuxfamily.org/dox/GettingStarted.html>
- Google & Internet

Part 1: k-NN for Classification

You will implement the k-NN algorithm for the classification problem in this task.

Dataset: The classification problem will be based on the **Iris dataset**, which is already available in the project folder.

The Visual Interface: The main form (i.e., [MainForm.h](#)) calls all the class objects and functions that are already implemented in the skeleton. However, you may need to edit some parts of the code to make it consistent with your code to display the results correctly.

- *Loading and splitting the dataset*

The code for loading and preparing the dataset is already implemented, so no additional code is required for this part. Use the following function in your code:

- You can use the function [loadAndPreprocessDataset\(\)](#) to prepare the dataset, which is located in the class name [DataLoader.cpp](#) under [DataUtils](#) folder.

- Use the `splitDataset()` function in the class name `DataPreprocessor.cpp` under the `DataUtils` folder to split the dataset into training and test sets.

Your Task: You will mainly work with two classes: `KNNClassifier.cpp` and `SimilarityFunctions.cpp`, located under the `Classification` and `Utils` folder, respectively.

- *Creating class member functions*

The `KNNClassifier.cpp` class may contain some or all of the following member functions:

- **Constructor:** initialises the member variables.
- **fit() function:** assigns training dataset and labels.
- **predict() function:** this is the main function that calculates the similarity between instances of test and training data sets and returns the classification prediction. *Note:* check the returns data type.

Task 1: In this Lab work you must complete the `predict()` function.

The `SimilarityFunctions.cpp` class contains several similarity functions bodies which are not implemented.

Task 2: You must complete `eulideanDistance()` function to compute the similarity between instances.

Task 3: In addition, you should at least implement one other distance function apart from the `eulideanDistance()` in the `SimilarityFunctions.cpp` class. Now, use that function for similarity measures in the `predict()` function. Does the algorithm perform the same when you use `eulideanDistance()`? Demonstrate and discuss the results with lab assistant.

- *Evaluating*

When you complete the implementation, it is time to evaluate your algorithm by measuring accuracy, precision, recall and F1 score. The code for these metrics is already done, so no additional implementation is required.

- You can use `accuracy()`, `precision()`, `recall()` and `f1Score()` functions, which are located in the class name `Metrics.cpp` under the `Evaluation` folder.

Task 4: What value of k have you used? Try different values for k . What is the largest and smallest value of k you can choose? Discuss your findings during the demonstration.

Note: The value of k is set when you create an object of the `KNNClassifier.cpp` class in the `MainForm.cpp` class.

- *Visualisation:*

A visual representation (for training and test data sets) is needed. The code for this task is already done, so use the following functions:

- Use the `confusionMatrix()` function located in the class name `Metrics.cpp` under the `Evaluation` folder, to calculate the confusion matrix for the specified number of class labels.
- Use the `DisplayConfusionMatrixInGrid()` function located in the `MainForm.h`, to Display the confusion matrix in the grid.

An example of the result from the k-NN classification should look shown in the Figure 1. The visual interface is same for all classifiers.

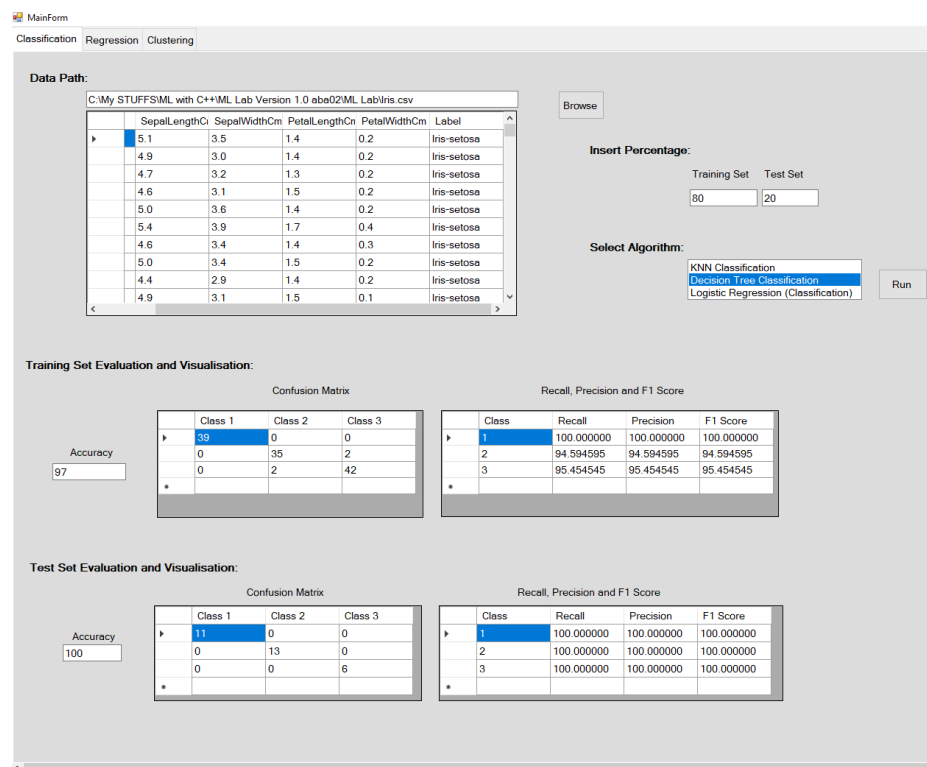


Figure 1: Example result of Decision Tree classification shown in the visual interface. The k-NN algorithm uses the same visual interface.

Part 2: k-NN for Regression

In this task, you will implement the k-NN algorithm for the regression problem.

Dataset

The regression problem will be based on the **Boston Housing** dataset, which is already available in the project folder.

The Visual Interface: Similar to the classification task, the main form (i.e., [MainForm.h](#)) calls all the class objects and functions that are already implemented. But you may need to edit some parts of the code to make it consistent with your code to display the results correctly.

Your Task: You will mainly work with the [KNNRegression.cpp](#) class located under the [Regression](#) folder. If you have already completed the **Part 1** then you do not need to implement anything more in the [SimilarityFunctions.cpp](#) class.

- *Loading and splitting the dataset*

The code for loading and preparing (preprocessing) the dataset is already done, so no additional code is required for this part. Use the following function in your code:

- You can use the function [readDatasetFromFilePath\(\)](#) to prepare the dataset, which is located in the class name [DataLoader.cpp](#) under [DataUtils](#) folder.
- Use the [splitDataset\(\)](#) function in the class name [DataPreprocessor.cpp](#) under the [DataUtils](#) folder to split the dataset into training and test sets.

- *Creating class member functions*

The [KNNRegression.cpp](#) class will have the following member functions:

- [Constructor](#): initialises the number of neighbours
- [fit\(\)](#) function: assigns training dataset and labels.
- [predict\(\)](#) function: this is the main function that calculates the similarity between instance of test and training data set and returns the prediction.
Note: check the returns data type.

Task 5: In this part of the work you must complete the [predict\(\)](#) function.

Task 6: If you have already completed **Task 2** and **Task 3**, you are good to go with the [SimilarityFunctions.cpp](#) class. Use the similarity function that you implemented in **Task 3** and show the result.

- Evaluating

Now, you need to evaluate the performance of the regression model measuring Mean Absolute Error (MAE), Root Mean Squared Error (RMSE) and R-Squared. The code for these metrics is already implemented.

- You can use `meanAbsoluteError()`, `rootMeanSquaredError()`, and `rSquared()` functions, which are located in the class name `Metrics` under the `Evaluation` folder for this evaluation.

Task 7: What value of k have you used? Try different values for k . What is the largest and smallest value of k you can choose? Discuss your findings during the demonstration.

Note: The value of k is set when you create an object of the `KNNRegression.cpp` class in the `MainForm.cpp` class.

- Visualisation

A visual representation (for training and test sets) is needed. The code for this task is already done, so use the following functions for that:

- Use the `VisualizeParityPlot()` function, located in the `MainForm.h`, to display the parity plot for the actual and predicted test labels.

An example of the result from the k -NN regression is shown in the Figure 2. The visual interface is same for all regression algorithms.

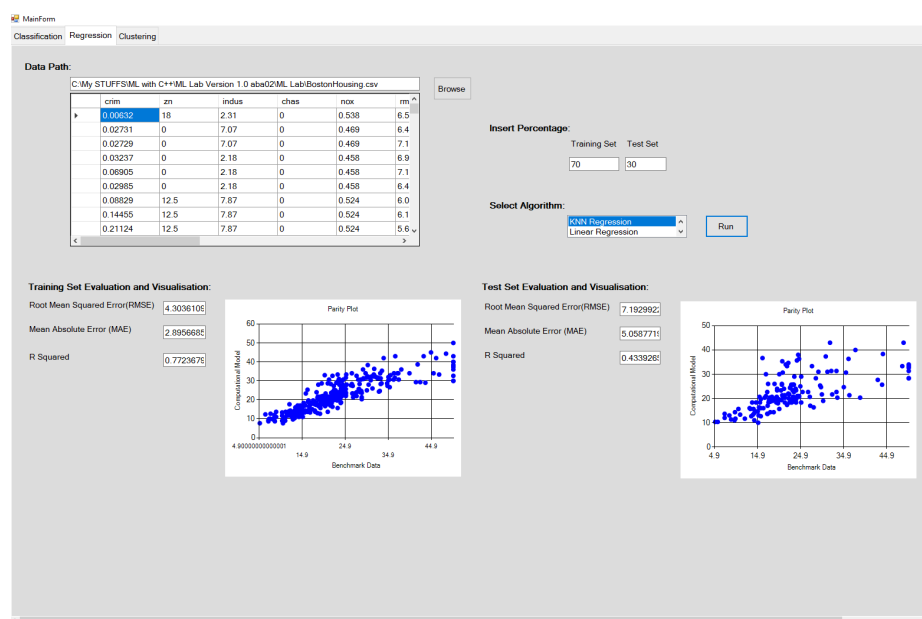


Figure 2: Example results of k -NN Regression in the visual interface.

Some Hints:

- Include the necessary header files at the beginning of your code.
- See lecture slides to predict labels considering k-nearest neighbours.
- Calculate the distance between the test data point and each training data point.
- Store predicted labels for all test data points.

Extra Assignment – if you want to learn some more.

Implement all the similarity functions in the [SimilarityFunctions.cpp](#) class. Test the performance of k-NN algorithm for both Classification and Regression using each similarity function in the [predict\(\)](#) function. Which similarity function works best?